Automatic Planning

7. Heuristic Search

How to Avoid Having to Look at a Gazillion States

Jörg Hoffmann



Winter Term 2017/2018

Jörg Hoffmann

Introduction

Automatic Planning

Chapter 7: Heuristic Search

Agenda

Introduction

- Introduction
- What Are Heuristic Functions?
- 3 How to Use Heuristic Functions?
- 4 How to Obtain Heuristic Functions?
- Conclusion

Jörg Hoffmann

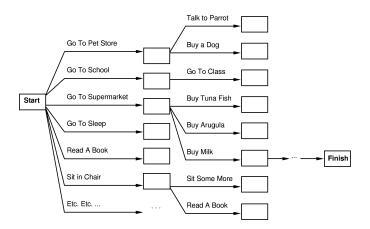
Fill in (some) details on these choices:

- Search space: Progression vs. regression.
 - ightarrow Previous Chapter
- Search algorithm: Uninformed vs. heuristic; systematic vs. local.
 - \rightarrow This Chapter
- Search control: Heuristic functions and pruning methods.
 - \rightarrow Chapters 8–19

t's a Heuristic? How to Use it? How to Obtain it? Conclusion References

Looking at a Gazillion States?

Introduction



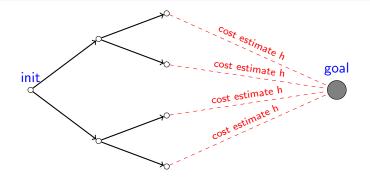
 \rightarrow Use heuristic function to guide the search towards the goal!

Jörg Hoffmann Automatic Planning Chapter 7: Heuristic Search 5/40

What's a Heuristic? How to Use it? How to Obtain it? Conclusion References

Heuristic Search

Introduction



 \rightarrow Heuristic function h estimates the cost of an optimal path from a state s to the goal; search prefers to expand states s with small h(s).

Live Demo vs. Breadth-First Search:

http://qiao.github.io/PathFinding.js/visual/

Jörg Hoffmann Automatic Planning Chapter 7: Heuristic Search 6/40

References

7/40

Our Agenda for This Chapter

Introduction

- What Are Heuristic Functions? Gives the basic definition, and introduces a number of important properties that we will be considering throughout the course.
- How to Use Heuristic Functions? Recaps the basic heuristic search algorithms from Al'17, and adds a few new ones. Gives a few planning-specific algorithms and explanations.
- 4 How to Obtain Heuristic Functions? Recaps the concept of "Relaxation" from Al'17: A basic explanation how heuristic functions are derived in practice.

Jörg Hoffmann Automatic Planning Chapter 7: Heuristic Search

Heuristic Functions

Definition (Heuristic Function). Let Π be a planning task with state space $\Theta_{\Pi} = (S, L, c, T, I, S^G)$. A heuristic function, short heuristic, for Π is a function $h: S \mapsto \mathbb{R}_0^+ \cup \{\infty\}$. Its value h(s) for a state s is referred to as the state's heuristic value, or h value.

Definition (Remaining Cost, h^*). Let Π be a planning task with state space $\Theta_{\Pi} = (S, L, c, T, I, S^G)$. For a state $s \in S$, the state's remaining cost is the cost of an optimal plan for s, or ∞ if there exists no plan for s. The perfect heuristic for Π , written h^* , assigns every $s \in S$ its remaining cost as the heuristic value.

- \rightarrow Heuristic functions h estimate remaining cost h^* .
- \rightarrow These definitions apply to both, STRIPS and FDR.

Jörg Hoffmann

Heuristic Functions: The Eternal Trade-Off

What does it mean, "estimate remaining cost"?

- In principle, the "estimate" is an arbitrary function. In practice, we want it to be accurate (aka: informative), i.e., close to the actual remaining cost.
- We also want it to be fast, i.e., a small overhead for computing h.
- These two wishes are in contradiction! Extreme cases?
 - $\rightarrow h = 0$: No overhead at all, completely un-informative. $h = h^*$: Perfectly accurate, overhead=solving the problem in the first place.
- \rightarrow We need to trade off the accuracy of h against the overhead of computing it. \rightarrow Chapters 8–13,15–17
- → What exactly is "accuracy"? How does it affect search performance? Interesting and challenging subject! We'll consider this in Chapter 18.

Questionnaire

Question!

For root-finding on a map, the straight-line distance heuristic certainly has small overhead. But is it accurate?

(A): No (B): Yes

(C): Sometimes (D): Maybe

- \rightarrow Depends on the map, and our initial location A and goal location B:
 - If there is a direct road from A to B, then straight-line distance is accurate (exact, in case the road has no curves at all).
 - If, say, A is central Africa and B is Patagonia, and we don't have boats capable of crossing an ocean, then the heuristic suggests to move to the African south-east coast while the actual solution is via Asia and North America . . .

Properties of Individual Heuristic Functions

Definition (Safe/Goal-Aware/Admissible/Consistent). Let Π be a planning task with state space $\Theta_{\Pi} = (S, L, c, T, I, S^G)$, and let h be a heuristic for Π . The heuristic is called:

- safe if, for all $s \in S$, $h(s) = \infty$ implies $h^*(s) = \infty$;
- goal-aware if h(s) = 0 for all goal states $s \in S^G$:
- admissible if $h(s) < h^*(s)$ for all $s \in S$;
- consistent if $h(s) \le h(s') + c(a)$ for all transitions $s \xrightarrow{a} s'$.

\rightarrow Relationships:

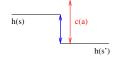
Proposition. Let Π be a planning task, and let h be a heuristic for Π . If h is admissible, then h is goal-aware. If h is admissible, then h is safe. If h is consistent and goal-aware, then h is admissible. No other implications of this form hold.

Proof. First two claims: Easy. Third claim: Next slide.

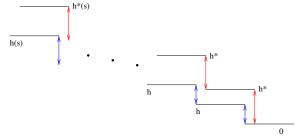
Consistency: Illustration

Introduction

Consistency = "heuristic value decrases by at most c(a)":



Consistent and goal-aware implies admissible: Let s be a state. $h^*(s)$ is the cost of an optimal solution path for s. Induction over that path, backwards from the goal: (on an optimal path, h^* decreases by exactly c(a) in each step)



Jörg Hoffmann

References

Properties of Individual Heuristic Functions, ctd.

Examples:

- Is h = Manhattan distance in the 15-Puzzle safe/goal-aware/admissible/consistent? All yes. Easy for goal-aware and safe (h is never ∞). Consistency: Moving a tile can't decrease h by more than 1.
- Is h =straight line distance safe/goal-aware/admissible/consistent? All yes. Easy for goal-aware and safe (h is never ∞). Consistency: If you drive 100km, then straight line distance can't decrease by more than 100km.
- An admissible but inconsistent heuristic: To-Moscow with h(SB)=1000, h(KL)=100.
- \rightarrow In practice, most heuristics are safe and goal-aware, and admissible heuristics are typically consistent.

What about inadmissible heuristics?

• Inadmissible heuristics typically arise as approximations of admissible heuristics that are too costly to compute. (Examples: Chapter 9)

Jörg Hoffmann Automatic Planning Chapter 7: Heuristic Search 14/40

Domination Between Heuristic Functions

Definition (Domination). Let Π be a planning task, and let h and h' be admissible heuristics for Π . We say that h' dominates h if $h \leq h'$, i.e., for all states s in Π we have $h(s) \leq h'(s)$.

 $\rightarrow h'$ dominates h= "h' provides a lower bound at least as good as h".

Remarks:

- Example: h' =Manhattan Distance vs. h =Misplaced Tiles in 15-Puzzle: Each misplaced tile accounts for at least 1 (typically, more) in h'.
- h^* dominates every other admissible heuristic.
- Modulo tie-breaking, the search space of A^* under h' can only be smaller than that under h. (See [Holte (2010)] for details)
- In Chapter 17, we will consider much more powerful concepts, comparing entire *families* of heuristic functions.

Jörg Hoffmann

Additivity of Heuristic Functions

Definition (Additivity). Let Π be a planning task, and let h_1, \ldots, h_n be admissible heuristics for Π . We say that h_1, \ldots, h_n are additive if $h_1 + \cdots + h_n$ is admissible, i.e., for all states s in Π we have $h_1(s) + \cdots + h_n(s) < h^*(s)$.

 \rightarrow An ensemble of heuristics is additive if its sum is admissible.

Remarks:

- Example: h_1 considers only tiles 1 ... 7, and h_2 considers only tiles 8 ... 15, in the 15-Puzzle: The two estimates are then, intuitively, "independent".
 - $(h_1 \text{ and } h_2 \text{ are orthogonal projections} \rightarrow \textbf{Chapter 12})$
- We can always combine h_1, \ldots, h_n admissibly by taking the max. Taking \sum is *much* stronger; in particular, \sum dominates \max .
- In Chapter 16, we will devise a third, strictly more general, technique to admissibly combine heuristic functions.

What Works Where in Planning?

Blind (no h) vs. heuristic:

- For satisficing planning, heuristic search vastly outperforms blind algorithms pretty much everywhwere.
- For optimal planning, heuristic search also is better (but the difference is not as huge).

Systematic (maintain all options) vs. local (maintain only a few) :

- For satisficing planning, there are successful instances of each.
- For optimal planning, systematic algorithms are required.
- \rightarrow Here, we briefly cover the search algorithms most successful in planning. For more details (in particular, for blind search), refer to Al'17 Chapters 4 and 5.

References

Reminder: Greedy Best-First Search and A*

For simplicity, duplicate elimination omitted and using Al'17 notation:

```
function Greedy Best-First Search [A*] (problem) returns a solution, or failure node \leftarrow a node n with n.state = problem.InitialState frontier notion \leftarrow a priority queue ordered by ascending n [g+h], only element n loop do if n [n if n if n
```

 \rightarrow Greedy best-first search explores states by increasing heuristic value h. A^* explores states by increasing plan-cost estimate g+h.

Jörg Hoffmann

Automatic Planning

Chapter 7: Heuristic Search

Greedy Best-First Search: Remarks

Properties:

Introduction

- Complete? Yes, with duplicate elimination. (If $h(s) = \infty$ states are pruned, h needs to be safe.)
- Optimal? No. (Even for perfect heuristics! E.g., say the start state has two transitions to goal states, one of which costs a million bucks while the other one is for free. Nothing keeps Greedy Best-First Search from choosing the bad one.)

Technicalities:

• Duplicate elimination: Insert child node n' only if n'. State is not already contained in *explored* \cup States(*frontier*). (Cf. Al'17)

Bottom line: Fast but not optimal \implies satisficing planning.

A*: Remarks

Properties:

- Complete? Yes. (Even without duplicate detection; if $h(s) = \infty$ states are pruned, h needs to be safe.)
- Optimal? Yes, for admissible heuristics.

Technicalities:

- "Plan-cost estimate" g(s) + h(s) known as f-value f(s) of s. \rightarrow If g(s) is taken from a cheapest path to s, then f(s) is a lower bound on the cost of a plan through s.
- Duplicate elimination: If n'.State \notin explored \cup States(frontier), then insert n'; else, insert n' only if the new path is cheaper than the old one, and if so remove the old path. (Cf. Al'17)

Bottom line: Optimal for admissible $h \implies$ optimal planning, with such h.

Weighted A*

Introduction

For simplicity, duplicate elimination omitted and using Al'17 notation:

```
function Weighted A*(problem) returns a solution, or failure
  node \leftarrow a \text{ node } n \text{ with } n.state = problem.InitialState
  frontier \leftarrow a priority queue ordered by ascending g + W *h, only element n
  loop do
       if Empty?(frontier) then return failure
       n \leftarrow Pop(frontier)
       if problem.GoalTest(n.State) then return Solution(n)
       for each action a in problem.Actions(n.State) do
          n' \leftarrow ChildNode(problem, n, a)
          Insert(n', [q(n') + W*h(n'), frontier)
```

 \rightarrow Weighted A* explores states by increasing weighted-plan-cost estimate q + W * h.

Jörg Hoffmann

Weighted A*: Remarks

The weight $W \in \mathbb{R}_0^+$ is an algorithm parameter:

- For W = 0, weighted A* behaves like? Uniform-cost search, i.e., "cheapest-first on path costs g".
- For W=1, weighted A^* behaves like? A^* .
- For $W=10^{100}$, weighted A^* behaves like? Greedy best-first search (i.e., if W is large enough, the "g" in "g+W*h" doesn't matter anymore.

Properties:

- For W > 1, weighted A^* is bounded suboptimal.
 - ightarrow If h is admissible, then the solutions returned are at most a factor W more costly than the optimal ones.

Bottom line: Allows to interpolate between greedy best-first search and A^* , trading off plan quality against computational effort.

Hill-Climbing

Introduction

Remarks:

- Is this complete or optimal? No.
- Can easily get stuck in local minima where immediate improvements of h(n) are not possible.
- Many variations: tie-breaking strategies, restarts, ... (cf. Al'17)

Jörg Hoffmann

References

```
function Enforced Hill-Climbing returns a solution node \leftarrow \text{a node } n \text{ with } n.state = problem.InitialState} \textbf{loop do} \textbf{if } problem.GoalTest(n.State) \textbf{ then return } Solution(n) \textbf{Perform breadth-first search for a node } n' \textbf{ s.t. } h(n') < h(n) n \leftarrow n'
```

Remarks:

- Is this optimal? No.
- Is this complete? See next slide.

Questionnaire

Introduction

```
\begin{array}{l} \textbf{function} \ \ \text{Enforced Hill-Climbing returns a solution} \\ node \leftarrow \ \ \text{a node } n \ \ \text{with } n.state = problem.InitialState} \\ \textbf{loop do} \\ \textbf{if } problem.GoalTest(n.State) \ \textbf{then return } Solution(n) \\ \textbf{Perform breadth-first search for a node } n' \ \text{s.t.} \ h(n') < h(n) \\ n \leftarrow n' \end{array}
```

Question!

Assume that h(s)=0 if and only if s is a goal state. Is Enforced Hill-Climbing complete?

- \rightarrow Only when restricting the input to planning tasks that do not contain any reachable unrecognized dead-end states:
 - ullet If there is a reachable unrecognized dead-end state, then the current node n may at some point end up containing that state, in which case the algorithm will not find a solution.
 - Say there are no reachable unrecognized dead-end states. Say the current node n contains the non-goal state s. Then h(s)>0, a goal state s' is reachable from s, and 0=h(s')< h(s). So breadth-first search will terminate with success.

Jörg Hoffmann Automatic Planning Chapter 7: Heuristic Search 26/40

hat's a Heuristic? How to Use it? How to Obtain it? Conclusion References

Heuristic Functions from Relaxed Problems

Introduction



Problem Π : Find a route from Saarbruecken To Edinburgh.

Jörg Hoffmann Automatic Planning Chapter 7: Heuristic Search 28/40

References

Heuristic Functions from Relaxed Problems

Introduction



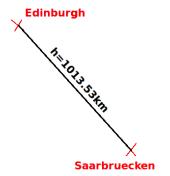


Relaxed Problem Π' : Throw away the map.

Jörg Hoffmann Automatic Planning Chapter 7: Heuristic Search 28/40

Heuristic Functions from Relaxed Problems

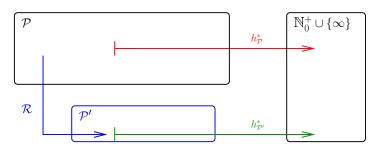
Introduction



Heuristic function h: Straight line distance.

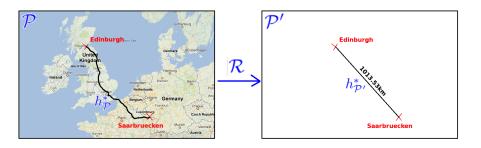
Jörg Hoffmann Automatic Planning Chapter 7: Heuristic Search 28/40

Introduction



- You have a class $\mathcal P$ of problems, whose perfect heuristic $h_{\mathcal P}^*$ you wish to estimate.
- You define a class \mathcal{P}' of *simpler problems*, whose perfect heuristic $h_{\mathcal{P}'}^*$ can be used to *estimate* $h_{\mathcal{P}}^*$.
- You define a transformation the relaxation mapping \mathcal{R} that maps instances $\Pi \in \mathcal{P}$ into instances $\Pi' \in \mathcal{P}'$.
- Given $\Pi \in \mathcal{P}$, you let $\Pi' := \mathcal{R}(\Pi)$, and estimate $h_{\mathcal{P}}^*(\Pi)$ by $h_{\mathcal{P}'}^*(\Pi')$.

Jörg Hoffmann



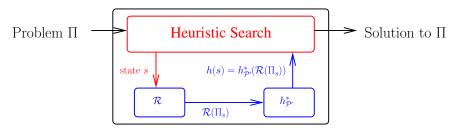
- Problem class \mathcal{P} : Route finding.
- Perfect heuristic $h_{\mathcal{D}}^*$ for \mathcal{P} : Length of a shortest route.
- Simpler problem class \mathcal{P}' : Route finding on an empty map.
- Perfect heuristic $h_{\mathcal{D}'}^*$ for \mathcal{P}' : Straight-line distance.
- Transformation \mathcal{R} : Throw away the map.

How to Relax During Search: Overview

Attention! Search uses the real (un-relaxed) Π . The relaxation is applied **only** within the call to h(s)!!!

How to Obtain it?

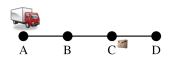
000000000



- Here, Π_s is Π with initial state replaced by s, i.e., $\Pi = (P, A, I, G)$ changed to (P, A, s, G): The task of finding a plan for search state s.
- A common student mistake is to instead apply the relaxation once to the whole problem, then doing the whole search "within the relaxation".
- Slides 34 and 32 illustrate the correct search process in detail.

Jörg Hoffmann

Chapter 7: Heuristic Search



Real problem:

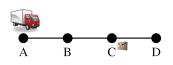
- Initial state *I*: *AC*; goal *G*: *AD*.
- Actions A: pre, add, del.
- \bullet drXY, loX, ulX.

Greedy best-first search:

Introduction

(tie-breaking: alphabetic)



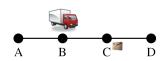


Relaxed problem:

- State s: AC; goal G: AD.
- Actions A: pre, add.
- $h^+(s) = 5$: e.g. $\langle drAB, drBC, drCD, loC, ulD \rangle$.

Greedy best-first search: (tie-breaking: alphabetic)



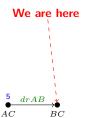


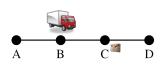
Real problem:

- State s: BC; goal G: AD.
- Actions A: pre, add, del.
- $AC \xrightarrow{drAB} BC$.

Greedy best-first search:

(tie-breaking: alphabetic)

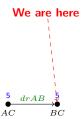




Relaxed problem:

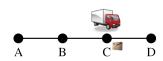
- State s: BC; goal G: AD.
- Actions A: pre, add.
- $h^+(s) = 5$: e.g. $\langle drBA, drBC, drCD, loC, ulD \rangle$.

Greedy best-first search: (tie-breaking: alphabetic)



Introduction

How to Relax During Search: Ignoring Deletes

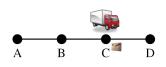


Real problem:

- State s: CC; goal G: AD.
- Actions A: pre, add, del.
- $BC \xrightarrow{drBC} CC$.

Greedy best-first search: (tie-breaking: alphabetic)



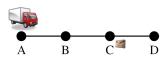


Relaxed problem:

- State s: CC; goal G: AD.
- Actions A: pre, add.
- $h^+(s) = 5$: e.g. $\langle drCB, drBA, drCD, loC, ulD \rangle$.

Greedy best-first search: (tie-breaking: alphabetic)

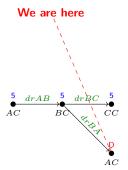




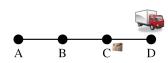
Real problem:

- State s: AC; goal G: AD.
- Actions A: pre, add, del.
- Duplicate state, prune.

Greedy best-first search: (tie-breaking: alphabetic)



How to Relax During Search: Ignoring Deletes

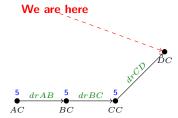


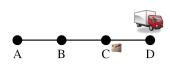
Real problem:

- State s: DC; goal G: AD.
- Actions A: pre, add, del.
- $CC \xrightarrow{drCD} DC$.

Greedy best-first search:

(tie-breaking: alphabetic)

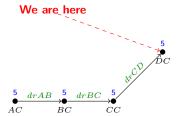




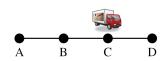
Relaxed problem:

- State s: DC; goal G: AD.
- Actions A: pre, add.
- $h^+(s) = 5$: e.g. $\langle drDC, drCB, drBA, loC, ulD \rangle$.

Greedy best-first search: (tie-breaking: alphabetic)



How to Relax During Search: Ignoring Deletes

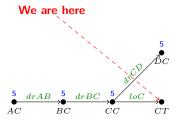


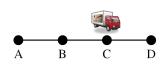
Real problem:

- State s: CT; goal G: AD.
- Actions A: pre, add, del.
- $CC \xrightarrow{loC} CT$.

Greedy best-first search:

(tie-breaking: alphabetic)

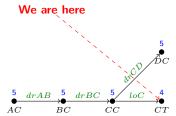




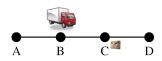
Relaxed problem:

- State s: CT; goal G: AD.
- Actions A: pre, add.
- $h^+(s) = 4$: e.g. $\langle drCB, drBA, drCD, ulD \rangle$.

Greedy best-first search: (tie-breaking: alphabetic)



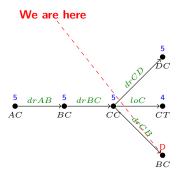
How to Relax During Search: Ignoring Deletes



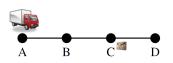
Real problem:

- State s: BC; goal G: AD.
- Actions A: pre, add, del.
- Duplicate state, prune.

Greedy best-first search: (tie-breaking: alphabetic)

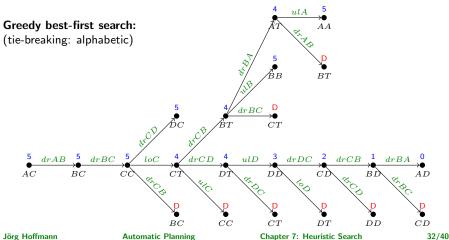


How to Relax During Search: Ignoring Deletes



Real problem:

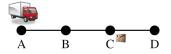
- Initial state *I*: *AC*; goal *G*: *AD*.
- Actions A: pre, add, del.
- $\bullet \ drXY, loX, ulX.$



A Simple Planning Relaxation: Only-Adds

Example: "Logistics"

Introduction

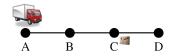


- Facts $P: \{truck(x) \mid x \in \{A, B, C, D\}\} \cup pack(x) \mid x \in \{A, B, C, D, T\}\}.$
- Initial state I: $\{truck(A), pack(C)\}$.
- Goal G: $\{truck(A), pack(D)\}.$
- Actions A: (Notated as "precondition \Rightarrow adds, \neg deletes")
 - drive(x, y), where x, y have a road: " $truck(x) \Rightarrow truck(y), \neg truck(x)$ ".
 - load(x): " $truck(x), pack(x) \Rightarrow pack(T), \neg pack(x)$ ".
 - unload(x): " $truck(x), pack(T) \Rightarrow pack(x), \neg pack(T)$ ".

Only-Adds Relaxation: Drop the preconditions and deletes.

"drive(x,y): $\Rightarrow truck(y)$ "; "load(x): $\Rightarrow pack(T)$ "; "unload(x): $\Rightarrow pack(x)$ ".

 \rightarrow Heuristic value for I is? 1: A plan for the relaxed task is $\langle unload(D) \rangle$.



Real problem:

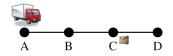
- Initial state *I*: *AC*; goal *G*: *AD*.
- Actions A: pre, add, del.
- \bullet drXY, loX, ulX.

Greedy best-first search:

Introduction

(tie-breaking: alphabetic)





Relaxed problem:

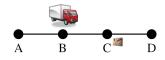
- State s: AC; goal G: AD.
- Actions A: add.
- $h^{\mathcal{R}}(s) = 1$: $\langle ulD \rangle$.

Greedy best-first search:

Introduction

(tie-breaking: alphabetic)



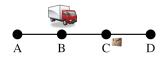


Real problem:

- State s: BC; goal G: AD.
- Actions A: pre, add, del.
- $AC \xrightarrow{drAB} BC$.

Greedy best-first search: (tie-breaking: alphabetic)





Relaxed problem:

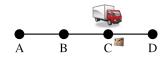
- State s: BC; goal G: AD.
- Actions A: add.
- $h^{\mathcal{R}}(s) = 2$: $\langle drBA, ulD \rangle$.

Greedy best-first search:

Introduction

(tie-breaking: alphabetic)



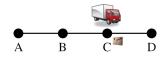


Real problem:

- State s: CC; goal G: AD.
- Actions A: pre, add, del.
- $BC \xrightarrow{drBC} CC$.

Greedy best-first search: (tie-breaking: alphabetic)



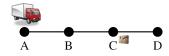


Relaxed problem:

- State s: CC; goal G: AD.
- Actions A: add.
- $h^{\mathcal{R}}(s) = 2$: $\langle drBA, ulD \rangle$.

Greedy best-first search: (tie-breaking: alphabetic)

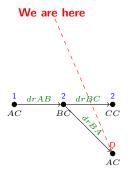
We are here AC

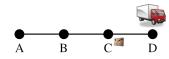


Real problem:

- State s: AC; goal G: AD.
- Actions A: pre, add, del.
- Duplicate state, prune.

Greedy best-first search: (tie-breaking: alphabetic)

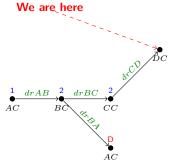


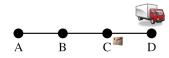


Real problem:

- State s: DC; goal G: AD.
- Actions A: pre, add, del.
- $CC \xrightarrow{drCD} DC$.

Greedy best-first search: (tie-breaking: alphabetic)

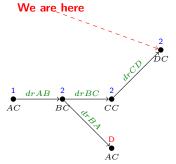


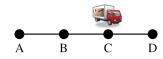


Relaxed problem:

- State s: DC; goal G: AD.
- Actions A: add.
- $h^{\mathcal{R}}(s) = 2$: $\langle drBA, ulD \rangle$.

Greedy best-first search: (tie-breaking: alphabetic)

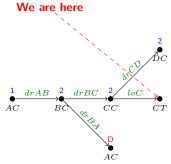


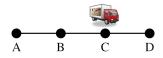


Real problem:

- State s: CT; goal G: AD.
- Actions A: pre, add, del.
- $CC \xrightarrow{loC} CT$.

Greedy best-first search: (tie-breaking: alphabetic)

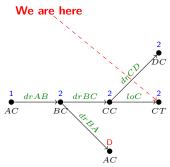


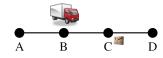


Relaxed problem:

- State s: CT; goal G: AD.
- Actions A: add.
- $h^{\mathcal{R}}(s) = 2$: $\langle drBA, ulD \rangle$.

Greedy best-first search: (tie-breaking: alphabetic)

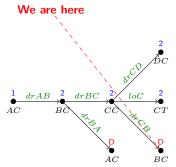


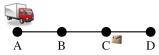


Real problem:

- State s: BC; goal G: AD.
- Actions A: pre, add, del.
- Duplicate state, prune.

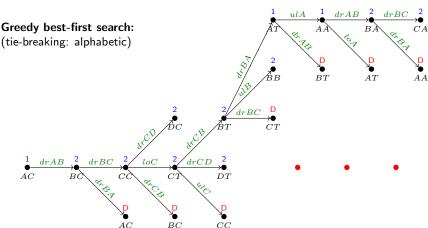
Greedy best-first search: (tie-breaking: alphabetic)





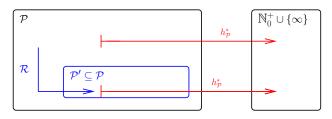
Real problem:

- Initial state I: AC; goal G: AD.
- Actions A: pre, add, del.
- \bullet drXY, loX, ulX.



Only-Adds and Ignoring Deletes are "Native" Relaxations

Native Relaxations: Confusing special case where $\mathcal{P}' \subseteq \mathcal{P}$.



- Problem class P: STRIPS planning tasks.
- Perfect heuristic $h_{\mathcal{P}}^*$ for \mathcal{P} : Length h^* of a shortest plan.
- ullet Transformation \mathcal{R} : Drop the (preconditions and) delete lists.
- Simpler problem class \mathcal{P}' is a special case of \mathcal{P} , $\mathcal{P}' \subseteq P$: STRIPS planning tasks with empty (preconditions and) delete lists.
- Perfect heuristic for \mathcal{P}' : Shortest plan for only-adds respectively delete-free STRIPS task.

Questionnaire

Question!

Is Only-Adds a "good heuristic" (accurate goal distance estimates) in ...

(A): Path Planning? (B): Blocksworld?

(C): Freecell? (D): SAT? (#unsatisfied clauses)

- ightarrow (A): No! The heuristic remains constantly 1 until we reach the actual goal state.
- \rightarrow (B): No: If we build a goal-tower of size 100 on top of a single block that still needs to move elsewhere, then the heuristic value is 1.
- \rightarrow (C): No: The heuristic value does take into account how many cards are already "home", but it is completely independent of the placement of all the other cards. In particular, dead-ends are essential in Freecell but the heuristic is completely unable to detect any of them.
- \rightarrow (D): No: Like in Freecell, the most essential part in SAT solving is knowing whether or not a given partial assignment is still feasible, i.e., whether or not it is a dead-end. The heuristic is completely unable to detect any of them.

Jörg Hoffmann Automatic Planning Chapter 7: Heuristic Search 36/40

Summary

Introduction

- Heuristic functions h map states to estimates of remaining cost. A heuristic
 can be safe, goal-aware, admissible, and/or consistent. A heuristic may
 dominate another heuristic, and an ensemble of heuristics may be additive.
- Greedy best-first search can be used for satisficing planning, A^* can be used for optimal planning provided h is admissible. Weighted A^* interpolates between the two.
- Relaxation is a method to compute heuristic functions. Given a problem \mathcal{P} we want to solve, we define a relaxed problem \mathcal{P}' . We derive the heuristic by mapping into \mathcal{P}' and taking the solution to this simpler problem as the heuristic estimate.
- During search, the relaxation is used *only inside* the computation of h(s) on each state s; the relaxation does not affect anything else.

References

Reading

Introduction

- Al'17 Chapters 4 and 5.
- A word of caution regarding Artificial Intelligence: A Modern Approach (Third Edition) [Russell and Norvig (2010)], Sections 3.6.2 and 3.6.3.

Content: These little sections are aimed at describing basically what I call "How to Relax" here. They do serve to get some intuitions. However, strictly speaking, they're a bit misleading. Formally, a pattern database (Section 3.6.3) is what is called a "relaxation" in Section 3.6.2: as we shall see in \rightarrow Chapters 11, 12, pattern databases are abstract transition systems that have more transitions than the original state space. On the other hand, not every relaxation can be usefully described this way; e.g., critical-path heuristics (\rightarrow Chapter 8) and ignoring-deletes heuristics (\rightarrow Chapter 9) are associated with very different state spaces.

References

References I

- Jörg Hoffmann and Bernhard Nebel. The FF planning system: Fast plan generation through heuristic search. *Journal of Artificial Intelligence Research*, 14:253–302, 2001.
- Robert C. Holte. Common misconceptions concerning heuristic search. In Ariel Felner and Nathan R. Sturtevant, editors, *Proceedings of the 3rd Annual Symposium on Combinatorial Search (SOCS'10)*, pages 46–51, Stone Mountain, Atlanta, GA, July 2010. AAAI Press.
- Stuart Russell and Peter Norvig. Artificial Intelligence: A Modern Approach (Third Edition). Prentice-Hall, Englewood Cliffs, NJ, 2010.